

CrashCoursesDL_FNN_CNN

January 10, 2024

#

Deep Learning Crash Course with tensorflow

Abdulrahman ALLOUCHE / 9 January 2024

Tensorflow is a free software Deep Learning library, developed by Google for the Python programming language.

<https://www.tensorflow.org/>

0.1 Set up TensorFlow

```
[57]: import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import warnings
warnings.filterwarnings("ignore")
print("TensorFlow version:", tf.__version__)
```

TensorFlow version: 2.9.0

0.2 1.Supervised DL using a Full Neural Networks

0.2.1 Load a dataset

```
[58]: mnist = tf.keras.datasets.mnist


(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
[59]: print("xtrain shape =", x_train.shape)
```

xtrain shape = (60000, 28, 28)

```
[60]: plt.figure(figsize=(3,3))
plt.imshow(x_test[0])
```

```
[60]: <matplotlib.image.AxesImage at 0x7f2618738220>
```



0.2.2 Build a model

```
[61]: warnings.filterwarnings("ignore")
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10)
])
model.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
flatten_4 (Flatten)	(None, 784)	0
dense_8 (Dense)	(None, 128)	100480
dropout_3 (Dropout)	(None, 128)	0
dense_9 (Dense)	(None, 10)	1290
Total params: 101,770		
Trainable params: 101,770		
Non-trainable params: 0		

```
[62]: #tf.keras.utils.plot_model(model, show_shapes=True) #need graphviz & pydot
```

```
[63]: predictions = model(x_train[:1]).numpy()
      predictions
```

```
[63]: array([[ -0.26856932, -0.28690484, -0.2719099 , -0.337744  ,  0.37312242,
          -0.45397654, -0.27955002,  0.21283133, -0.33580464,  0.00139271]],
      dtype=float32)
```

The `tf.nn.softmax` function converts these logits to probabilities for each class:

```
[64]: tf.nn.softmax(predictions).numpy()
```

```
[64]: array([[0.08699868, 0.08541805, 0.08670855, 0.08118401, 0.1652707 ,
          0.07227554, 0.08604861, 0.14079341, 0.08134161, 0.11396085]],
      dtype=float32)
```

`losses.SparseCategoricalCrossentropy` takes a vector of logits and a True index and returns a scalar loss for each example. This loss is equal to the negative log probability of the true class: **The loss is zero if the model is sure of the correct class.**

```
[65]: loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)
```

```
[66]: #tf.keras.losses. # tab key to obtain the list of available fonctions
```

```
[67]: loss_fn(y_train[:1], predictions).numpy()
```

```
[67]: 2.6272695
```

0.2.3 Before you start training, configure and compile the model

```
[68]: model.compile(optimizer='adam',
                  loss=loss_fn,
                  metrics=['accuracy'])
```

0.2.4 Train and evaluate your model

```
[69]: model.fit(x_train, y_train, epochs=5)
      #print(x_train.shape[0]/32)
      #model.fit(x_train, y_train, epochs=5, validation_split=0.2)
```

Epoch 1/5

1875/1875 [=====] - 3s 1ms/step - loss: 0.2957 - accuracy: 0.9137

Epoch 2/5

1875/1875 [=====] - 2s 1ms/step - loss: 0.1430 - accuracy: 0.9575

Epoch 3/5

1875/1875 [=====] - 2s 1ms/step - loss: 0.1065 - accuracy: 0.9677

Epoch 4/5

```
1875/1875 [=====] - 2s 1ms/step - loss: 0.0871 -
accuracy: 0.9728
Epoch 5/5
1875/1875 [=====] - 2s 1ms/step - loss: 0.0758 -
accuracy: 0.9766
```

[69]: <keras.callbacks.History at 0x7f26186ef400>

```
[70]: #callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=2)
callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=3)
history = model.fit(x_train, y_train, epochs=20, callbacks=[callback],
    ↪ validation_split=0.2)
len(history.history['loss'])
```

```
Epoch 1/20
1500/1500 [=====] - 2s 2ms/step - loss: 0.0658 -
accuracy: 0.9788 - val_loss: 0.0414 - val_accuracy: 0.9868
Epoch 2/20
1500/1500 [=====] - 2s 2ms/step - loss: 0.0573 -
accuracy: 0.9811 - val_loss: 0.0424 - val_accuracy: 0.9858
Epoch 3/20
1500/1500 [=====] - 2s 1ms/step - loss: 0.0505 -
accuracy: 0.9834 - val_loss: 0.0435 - val_accuracy: 0.9858
Epoch 4/20
1500/1500 [=====] - 2s 2ms/step - loss: 0.0459 -
accuracy: 0.9845 - val_loss: 0.0458 - val_accuracy: 0.9846
```

[70]: 4

The Model.evaluate method checks the models performance, usually on a “Validation-set” or “Test-set”.

```
[71]: model.evaluate(x_test, y_test, verbose=2)
```

```
313/313 - 0s - loss: 0.0694 - accuracy: 0.9801 - 323ms/epoch - 1ms/step
```

[71]: [0.0694371908903122, 0.9800999760627747]

If you want your model to return a probability, you can wrap the trained model, and attach the softmax to it:

```
[72]: probability_model = tf.keras.Sequential([
    model,
    tf.keras.layers.Softmax()
])
```

```
[73]: probability_model(x_test[:1])
```

[73]: <tf.Tensor: shape=(1, 10), dtype=float32, numpy=
array([[1.4111662e-09, 5.9732941e-10, 5.4968194e-08, 4.8179479e-04,


```
1.9729827e-14, 8.8700549e-09, 3.1283237e-14, 9.9951530e-01,  
1.2299834e-07, 2.7611291e-06]], dtype=float32)>
```

```
[74]: #Print predicted values vs real values  
def printErrValues(x_test,y_test, model):  
    probability_model = tf.keras.Sequential([  
        model,  
        tf.keras.layers.Softmax()  
    ])  
    class_names = [0, 1, 2, 3, 4,5, 6, 7, 8, 9]  
    #n=1000  
    n=y_test.shape[0]  
    class_namesp = tf.constant([class_names]*n)  
    class_namesp = tf.reshape(class_namesp, [n*10])  
    pm=probability_model(x_test[0:n])  
    t=(pm>=tf.reduce_max(pm))  
    #print(t)  
    t = tf.reshape(t, [n*10])  
    #print(t.shape)  
    #print(class_namesp.shape)  
    #print(class_namesp[t].shape)  
  
    #print(class_namesp[t].numpy())  
    #print(y_test[0:n+1])  
    y_pred= class_namesp[t].numpy()  
    nerr=0  
    for i in range(y_pred.shape[0]):  
        if y_test[i] != y_pred[i]:  
            nerr += 1  
            #print("#index = ",i,' y_test=',y_test[i], ' y_predic=',y_pred[i])  
    print("nerr/n=",nerr, " / ",n)
```

```
[75]: #Print predicted values vs real values  
printErrValues(x_test,y_test, model)
```


```
nerr/n= 583 / 10000
```

```
[76]: plt.figure()  
plt.imshow(x_test[0])  
plt.colorbar()  
plt.grid(False)  
plt.show()
```



```
[77]: class_names = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']
```

```
plt.figure(figsize=(10,10))
for i in range(5):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(x_test[i], cmap=plt.cm.binary)
    plt.xlabel(class_names[y_test[i]])
plt.show()
```



0.2.5 Build a model with Functional API

```
[78]: inputs = tf.keras.Input(shape=(28,28))
      flatten= tf.keras.layers.Flatten(input_shape=(28, 28))
      x=flatten(inputs)
      dense = tf.keras.layers.Dense(128, activation="relu")
      x = dense(x)
      dropout=tf.keras.layers.Dropout(0.2)
      x = dropout(x)
      outputs = tf.keras.layers.Dense(10)(x)
      model = tf.keras.Model(inputs=inputs, outputs=outputs, name="mnist_model")
```

```
[79]: model.summary()
```

Model: "mnist_model"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 28, 28)]	0
flatten_5 (Flatten)	(None, 784)	0
dense_10 (Dense)	(None, 128)	100480
dropout_4 (Dropout)	(None, 128)	0
dense_11 (Dense)	(None, 10)	1290

=====
Total params: 101,770
Trainable params: 101,770
Non-trainable params: 0
=====

```
[80]: #tf.keras.utils.plot_model(model)
      #tf.keras.utils.plot_model(model, "my_first_model_with_shape_info.png",
      ↪show_shapes=True)
      tf.keras.utils.plot_model(model, show_shapes=True)
```

You must install pydot (``pip install pydot``) and install graphviz (see instructions at <https://graphviz.gitlab.io/download/>) for `plot_model/model_to_dot` to work.

```
[81]: model.compile(
      loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
      optimizer=tf.keras.optimizers.Adam(),
      metrics=["accuracy"],
      )
```

```
[82]: #reload data if necessary
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
[83]: history = model.fit(x_train, y_train, batch_size=64, epochs=5,
    ↪validation_split=0.2)

test_scores = model.evaluate(x_test, y_test, verbose=2)
print("Test loss:", test_scores[0])
print("Test accuracy:", test_scores[1])
```

```
Epoch 1/5
750/750 [=====] - 2s 2ms/step - loss: 0.3765 -
accuracy: 0.8931 - val_loss: 0.1729 - val_accuracy: 0.9534
Epoch 2/5
750/750 [=====] - 1s 2ms/step - loss: 0.1788 -
accuracy: 0.9481 - val_loss: 0.1334 - val_accuracy: 0.9621
Epoch 3/5
750/750 [=====] - 1s 2ms/step - loss: 0.1318 -
accuracy: 0.9614 - val_loss: 0.1087 - val_accuracy: 0.9677
Epoch 4/5
750/750 [=====] - 1s 2ms/step - loss: 0.1091 -
accuracy: 0.9672 - val_loss: 0.0956 - val_accuracy: 0.9722
Epoch 5/5
750/750 [=====] - 1s 2ms/step - loss: 0.0905 -
accuracy: 0.9732 - val_loss: 0.0888 - val_accuracy: 0.9737
313/313 - 0s - loss: 0.0827 - accuracy: 0.9742 - 238ms/epoch - 762us/step
Test loss: 0.08268717676401138
Test accuracy: 0.9742000102996826
```

0.3 2. Convolutional Neural Network (CNN)

0.3.1 Load a dataset

```
[84]: mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

x_train = x_train.reshape((x_train.shape[0], 28, 28, 1))
x_test = x_test.reshape((x_test.shape[0], 28, 28, 1))
```


0.3.2 Build a model

```
[85]: model = tf.keras.models.Sequential()

model.add(tf.keras.layers.Conv2D(28, (3, 3), activation='relu',
    ↪input_shape=(28, 28,1)))
model.add(tf.keras.layers.MaxPooling2D((2, 2)))
model.add(tf.keras.layers.Conv2D(56, (3, 3), activation='relu'))
model.add(tf.keras.layers.MaxPooling2D((2, 2)))
model.add(tf.keras.layers.Conv2D(56, (3, 3), activation='relu'))
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(56, activation='relu'))
model.add(tf.keras.layers.Dense(10))
```

```
[86]: model.summary()
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 26, 26, 28)	280
max_pooling2d_2 (MaxPooling 2D)	(None, 13, 13, 28)	0
conv2d_4 (Conv2D)	(None, 11, 11, 56)	14168
max_pooling2d_3 (MaxPooling 2D)	(None, 5, 5, 56)	0
conv2d_5 (Conv2D)	(None, 3, 3, 56)	28280
flatten_6 (Flatten)	(None, 504)	0
dense_12 (Dense)	(None, 56)	28280
dense_13 (Dense)	(None, 10)	570
Total params: 71,578		
Trainable params: 71,578		
Non-trainable params: 0		

```
[87]: #tf.keras.utils.plot_model(model, show_shapes=True)
```

0.3.3 Compile the model

```
[88]: model.compile(optimizer='adam',  
                  loss=tf.keras.losses.  
                  ↪SparseCategoricalCrossentropy(from_logits=True),  
                  metrics=['accuracy'])
```

0.3.4 Train and evaluate your model

```
[89]: model.fit(x_train, y_train, epochs=5)
```

```
Epoch 1/5  
1875/1875 [=====] - 13s 7ms/step - loss: 0.1539 -  
accuracy: 0.9517  
Epoch 2/5  
1875/1875 [=====] - 12s 7ms/step - loss: 0.0463 -  
accuracy: 0.9858  
Epoch 3/5  
1875/1875 [=====] - 13s 7ms/step - loss: 0.0339 -  
accuracy: 0.9893  
Epoch 4/5  
1875/1875 [=====] - 15s 8ms/step - loss: 0.0261 -  
accuracy: 0.9917  
Epoch 5/5  
1875/1875 [=====] - 14s 8ms/step - loss: 0.0207 -  
accuracy: 0.9936
```

```
[89]: <keras.callbacks.History at 0x7f26182b7ac0>
```

```
[90]: model.evaluate(x_test, y_test, verbose=2)
```

```
313/313 - 1s - loss: 0.0351 - accuracy: 0.9895 - 753ms/epoch - 2ms/step
```

```
[90]: [0.0351063534617424, 0.9894999861717224]
```

```
[91]: printErrValues(x_test,y_test, model)
```

```
nerr/n= 885 / 10000
```

0.4 Metrics

```
[92]: from sklearn.metrics import confusion_matrix, accuracy_score, precision_score,  
      ↪recall_score, f1_score  
      import seaborn as sns
```

```
[93]: probability_model = tf.keras.Sequential([  
    model,  
    tf.keras.layers.Softmax()  
)
```

```

prob=probability_model(x_test)
print(prob.shape)
y_pred=tf.math.argmax(prob,axis=1)
print(y_pred.shape)
print(y_test.shape)
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True,fmt='d')


```

(10000, 10)

(10000,)

(10000,)

[93]: <AxesSubplot:>



```

[94]: accuracy = accuracy_score(y_test, y_pred)
'''
precision = precision_score(y_test, y_pred,average=None)
recall = recall_score(y_test, y_pred,average=None)
f1score = f1_score(y_test, y_pred,average=None)
'''
precision = precision_score(y_test, y_pred,average="weighted")
recall = recall_score(y_test, y_pred,average="weighted")
f1score = f1_score(y_test, y_pred,average="weighted")

print(f"Accuracy = {accuracy}")
print(f"Precision = {precision}")

```

```
print(f"Recall = {recall}")
print(f"F1 Score = {f1score}")
```

Accuracy = 0.9895
Precision = 0.9895755607403444
Recall = 0.9895
F1 Score = 0.9895035541756534

```
[95]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	980
1	1.00	0.99	0.99	1135
2	0.99	0.99	0.99	1032
3	0.99	0.99	0.99	1010
4	1.00	0.99	0.99	982
5	0.97	0.99	0.98	892
6	0.99	0.99	0.99	958
7	1.00	0.99	0.99	1028
8	0.99	0.99	0.99	974
9	0.99	0.97	0.98	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

0.5 3. Choose devices

```
[96]: print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
gpus = tf.config.list_physical_devices('GPU')
gpus
```

Num GPUs Available: 0 GPU compatible avec tensorflow

```
[96]: []
```

```
[97]: print("Num CPUs Available: ", len(tf.config.list_physical_devices('CPU')))
cpus = tf.config.list_physical_devices('CPU')
cpus
```

Num CPUs Available: 1

```
[97]: [PhysicalDevice(name='/physical_device:CPU:0', device_type='CPU')]
```

```
[98]: with tf.device('/CPU:0'):
# Create some tensors
```

choice of using a given number of GPUs and CPUs

```
a = tf.constant([[1.0, 2.0, 3.0], [4.0, 5.0, 6.0]])
b = tf.constant([[1.0, 2.0], [3.0, 4.0], [5.0, 6.0]])
c = tf.matmul(a, b)
print(c)
```

```
tf.Tensor(
[[22. 28.]
 [49. 64.]], shape=(2, 2), dtype=float32)
```

```
[99]: tf.config.set_visible_devices(cpus[0], 'CPU')
print(1800*3)
```

```
5400
```

```
[100]: tf.config.set_visible_devices([], 'GPU') # deactivate GPU
```

```
[ ]:
```

```
[ ]:
```